

Text Data in Business and Economics

University of Basel, 21–25 September 2026

Benjamin W. Arold, University of Cambridge

11. Audio as Data

Outline

Introduction

Classical Machine Learning

Recurrent Neural Nets

Human Hearing is Interesting

- ▶ Hearing is a critical sense for communication and survival
- ▶ Like vision, it enables us to interact with the world without physical contact
- ▶ Various areas of the brain are associated with auditory processing
- ▶ Humans are "**analysts**" and "**generators**" of audio at the same time
- ▶ Human hearing has constraints:
 - ▶ Limited frequency range: 20 Hz to 20 kHz
 - ▶ Difficulty in locating low-frequency sounds
- ▶ Fascinating questions:
 - ▶ Fundamental: How does the auditory system decode sounds? (Beyond scope of this class)
 - ▶ Aggregate/social: What is the impact of sound on societal outcomes? (Central to our class)

Audio Analysis Definitions

- ▶ Many view it as a signal-processing challenge
- ▶ Julius O. Smith III: reconstructing auditory experiences
- ▶ This class explores techniques to analyze auditory data to study economics questions
- ▶ Focus on existing auditory data, not the mechanics of data capture

What is Sound?

- ▶ A sensory experience created by vibrations traveling through a medium (usually air)
- ▶ Analogy to ocean waves: sound waves propagate through air like ocean waves move through water
- ▶ As ocean waves interact with air, they create ripples in the atmosphere, analogous to **sound waves**




Figure: Ocean waves analogy to sound waves (Source: Sulzer, 2021)

Humans typically hear between 20 Hz and 20 kHz.

- ▶ We begin to hear low vibrations at around 20 Hz
- ▶ Teenagers can perceive frequencies up to about 20,000 Hz.
- ▶ Each octave doubles the frequency.
- ▶ Humans with good hearing can perceive a ten-octave range.

1. 20 Hz
2. 40 Hz ($20 \text{ Hz} \times 2$)
3. 80 Hz ($40 \text{ Hz} \times 2$)
4. 160 Hz ($80 \text{ Hz} \times 2$)
5. 320 Hz ($160 \text{ Hz} \times 2$)
6. 640 Hz ($320 \text{ Hz} \times 2$)
7. 1280 Hz ($640 \text{ Hz} \times 2$)
8. 2560 Hz ($1280 \text{ Hz} \times 2$)
9. 5120 Hz ($2560 \text{ Hz} \times 2$)
10. 10240 Hz ($5120 \text{ Hz} \times 2$)
11. 20480 Hz ($10240 \text{ Hz} \times 2$)

Figure: Human hearing range

Key Concepts: Hz, Octave, and Sampling Frequency

- ▶ **Hz (Hertz):** Measures the number of cycles per second of a periodic waveform
 - ▶ Higher Hertz means a higher-pitched sound
 - ▶ If a tone has a frequency of 440 Hz, the sound wave completes 440 cycles every second (A4)
- ▶ **Octave:** The interval between one musical pitch and another with half or double its frequency
 - ▶ 440 Hz and 880 Hz
- ▶ **Sampling frequency:** The number of samples per second from a continuous signal to make a digital signal:
 - ▶ Standard CD sampling rate is 44.1 kHz...
 - ▶ ...meaning that the audio waveform is being sampled 44,100 times per second

The wavelength is an important sound wave characteristic

- ▶ What is a sound wave, to begin with?
 - ▶ A transfer of sound energy through a medium
 - ▶ It is transmitted via the vibration of particles within that medium
- ▶ Illustration for non-physicists:
 - ▶ When something vibrates (like a speaker's diaphragm), it pushes on the neighboring air particles, increasing the pressure in that region → this push is then transferred from particle to particle → when it reaches someone's ear, the ear interprets the varying pressures as sound
- ▶ The physical distance between identical points in consecutive cycles of a sound wave
 - ▶ Determines the pitch

Frequency and Wavelength

The wavelength of a sound wave depends on its frequency.

$$\lambda = \frac{v}{f}$$

- ▶ λ = wavelength (distance between two wave peaks)
- ▶ v = speed of sound in air (≈ 343 m/s)
- ▶ f = frequency in Hertz

Implication

- ▶ Higher frequency $\rightarrow$ shorter wavelength
- ▶ Lower frequency $\rightarrow$ longer wavelength

Example

- ▶ A tone of 440 Hz (concert A) has wavelength

$$\lambda = \frac{343}{440} \approx 0.78 \text{ meters}$$

The amplitude is another important characteristic

- ▶ Amplitude in a general wave context: the magnitude of change in the oscillating variable within the wave
 - ▶ → Larger amplitudes mean more energy and often translate to louder sounds when talking about sound waves
- ▶ Amplitude in digital audio signal processing : understood as instantaneous amplitudes that represent the “magnitude of change” in the pressure wave (the sound wave) from its equilibrium position at that particular moment
 - ▶ The specific value of the audio waveform at a given sample point

One-Tone Sound Wave




Figure: Sidney Wood and SWPhonetics, 1994-2024

A More Complex Sound Wave




Figure: Sidney Wood and SWPhonetics, 1994-2024

Outline

Introduction

Classical Machine Learning

Recurrent Neural Nets

Recap from computer vision: For classical ML, we (often) extract features explicitly.

- Recall the typical pipeline for classical machine learning:




Figure: Dey (2018)¹

¹Dey, S. (2018). Hands-On Image Processing with Python: Expert Techniques for Advanced Image analysis and Effective Interpretation of Image Data. Packt Publishing Ltd.

We can use a similar workflow for audio data.

- Recall the typical pipeline for classical machine learning:




Figure: Adaptation by Widmer (2023) of Dey (2018)²

²Dey, S. (2018). Hands-On Image Processing with Python: Expert Techniques for Advanced Image analysis and Effective Interpretation of Image Data. Packt Publishing Ltd.

Can we use the raw audio input for classification tasks?

- ▶ Neural networks, such as CNN, can handle raw audio data → extract relevant features automatically (cf. images part)
- ▶ For very simple tasks, it may be possible with classical machine learning
- ▶ More generally, audio waveforms are high-dimensional data, especially for long clips
 - ▶ Computationally demanding
 - ▶ And often not necessary → suitable format depends on task
- ▶ For both classical and neural approaches, feature extraction is often useful

Spectral features capture important characteristics of a sound's frequency content

- ▶ Spectral features are widely used for audio tasks (e.g., music analysis, speech processing, audio classification)
- ▶ Typically, they encompass spectral centroid, spectral bandwidth, or spectral flatness
- ▶ The spectral centroid shows the frequency spectrum's “center of gravity”
 - ▶ Calculated as the weighted mean of the frequencies in the signal, with their amplitudes being the weights
 - ▶ Higher values reflect brighter sound
 - ▶ Technically, we cut the sound into frames and apply a Fourier Transform per frame
 - ▶ Hence, we transform the frame's time-domain signal into the frequency-domain

Spectral features beyond the spectral centroid

- ▶ The spectral bandwidth describes how wide the frequency band is
 - ▶ A wider bandwidth implies a broad range of frequencies
 - ▶ Conversely, a narrower bandwidth indicates a more tonally pure sound
 - ▶ It can help, for example, to identify the complexity of sound
- ▶ The spectral flatness indicates how “noise-like” a sound is instead of being tonal
 - ▶ Values close to 1 indicate a noise-like sound
 - ▶ Values near 0 indicate a more tonal sound
 - ▶ It can be helpful, for instance, to disentangle tones and environmental noises

Mel-Frequency Cepstral Coefficients (MFCCs) are often used for feature extraction

- ▶ MFCCs summarize the spectral structure of a sound in a way that reflects human hearing.
- ▶ Humans perceive sound frequencies non-linearly (rather, logarithmically)
- ▶ Human sensitivity is greater to changes in lower frequencies
- ▶ The Mel scale is a way to measure pitch that matches how we hear sounds
 - ▶ It is a perceptual scale of pitches judged to be equal in distance from one another
 - ▶ Originally derived from experiments with human listeners
- ▶ MFCCs represent the power spectrum of an audio signal more in line with human hearing
 - ▶ They use the Mel scale

MFCCs capture the short-term power spectrum of sound

- ▶ One typically begins by dividing the audio signal into short (overlapping) frames, for instance 20-40 ms
 - ▶ This allows assuming stationarity within each frame
 - ▶ Stationarity means that the statistical properties of the signal (like mean, variance) are constant over the frame's duration
- ▶ Several processing steps involved (including fast fourier transform)
- ▶ Bottom line: Typically end up with 12-13 coefficients per frame
- ▶ These coefficients capture the **overall spectral envelope of the sound**.
 - ▶ Empirically found to capture the most important features

Outline

Introduction

Classical Machine Learning

Recurrent Neural Nets

Recall that feature extraction is used for many audio (classification) tasks

- Recall the typical pipeline for classical machine learning




Figure: Adaptation by Widmer (2023) of Dey (2018)³

³Dey, S. (2018). Hands-On Image Processing with Python: Expert Techniques for Advanced Image analysis and Effective Interpretation of Image Data. Packt Publishing Ltd.

What models are used in audio analysis with deep learning?

- ▶ The audio analysis landscape is diverse
- ▶ CNN are used for classification with audio data
 - ▶ But less prominently than in CV
 - ▶ Audio data is often transformed into spectrograms (time-frequency representations) to use CNN
 - ▶ Examples: Emotions, speaker identity, accents
- ▶ Otherwise, varied, possibly hybrid approaches (RNN, transformers)
- ▶ In any case, audio data often requires pre-processing
- ▶ Other methods in audio analysis:
 - ▶ Classical ML techniques (e.g., MFCC)
 - ▶ Unsupervised and semi-supervised approaches: Useful with limited labelled data

Audio data is sometimes used with CNN or RNN

- ▶ Tailored for tasks with temporal dependencies
- ▶ For CNN, the conceptual foundation discussed in the computer vision part also applies here
- ▶ Recurrent Neural Networks (RNN) are “specialists” for sequential data
 - ▶ Sequential data examples are audio or text
 - ▶ RNN have a memory mechanism to remember previous inputs
 - ▶ They can, thus, “remember” context in sequences
 - ▶ For sequences, context is important to understand the overall pattern

Conceptually, how does an RNN proceed

- ▶ The data is typically pre-processed
 - ▶ Extract features: for example, MFCC by time t (often relatively short snippets)
 - ▶ Each feature at time t becomes input x_t to the RNN
- ▶ $\forall t \in \{0, 1, 2, \dots, T\}$, x_t are processed by the RNN sequentially
- ▶ RNN computes new hidden state h_t for each time step t
- ▶ Typically: $h_t = \text{Activation}(W_{xh} \cdot x_t + W_{hh} \cdot h_{t-1} + b)$
- ▶ W_{xh} , W_{hh} are weight matrices; b is a bias vector
- ▶ RNN learns W_{xh} , W_{hh} , and b
- ▶ Hidden state h_t used for output (e.g., in classification)